

Learning fragments of the TCP communication protocol

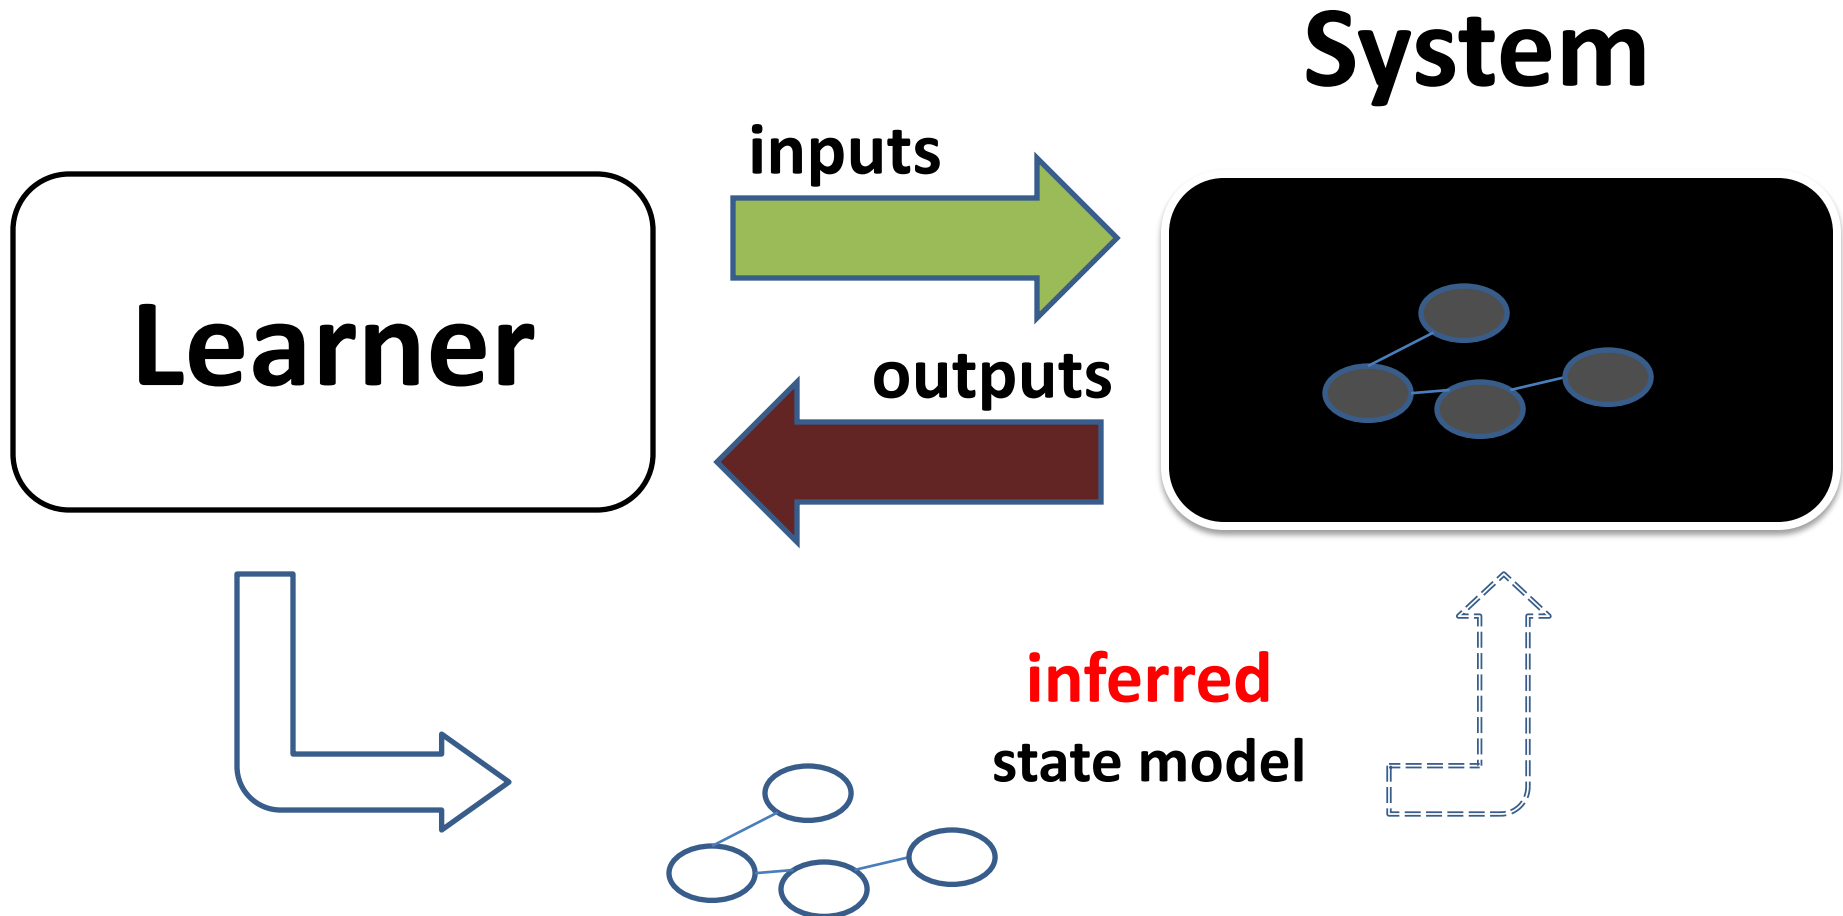
using **abstraction**

Radboud University



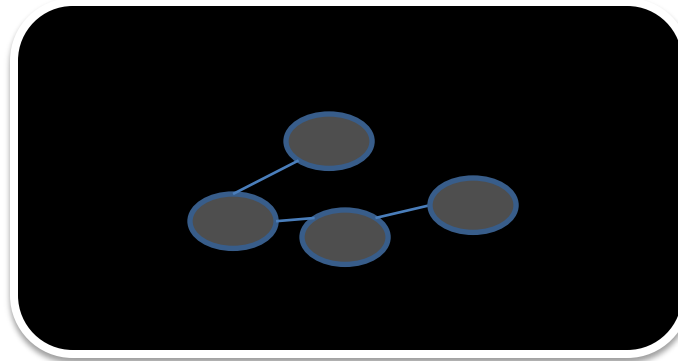
Paul Fiterau,
Frits Vaandrager, Ramon Janssen

Introduction

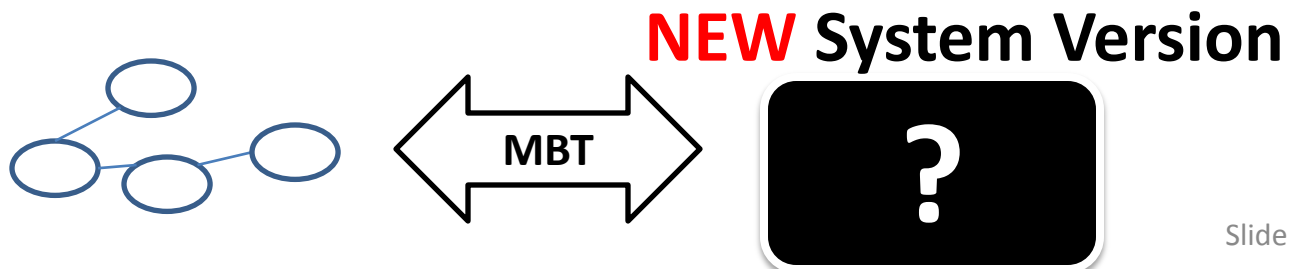


Introduction

System



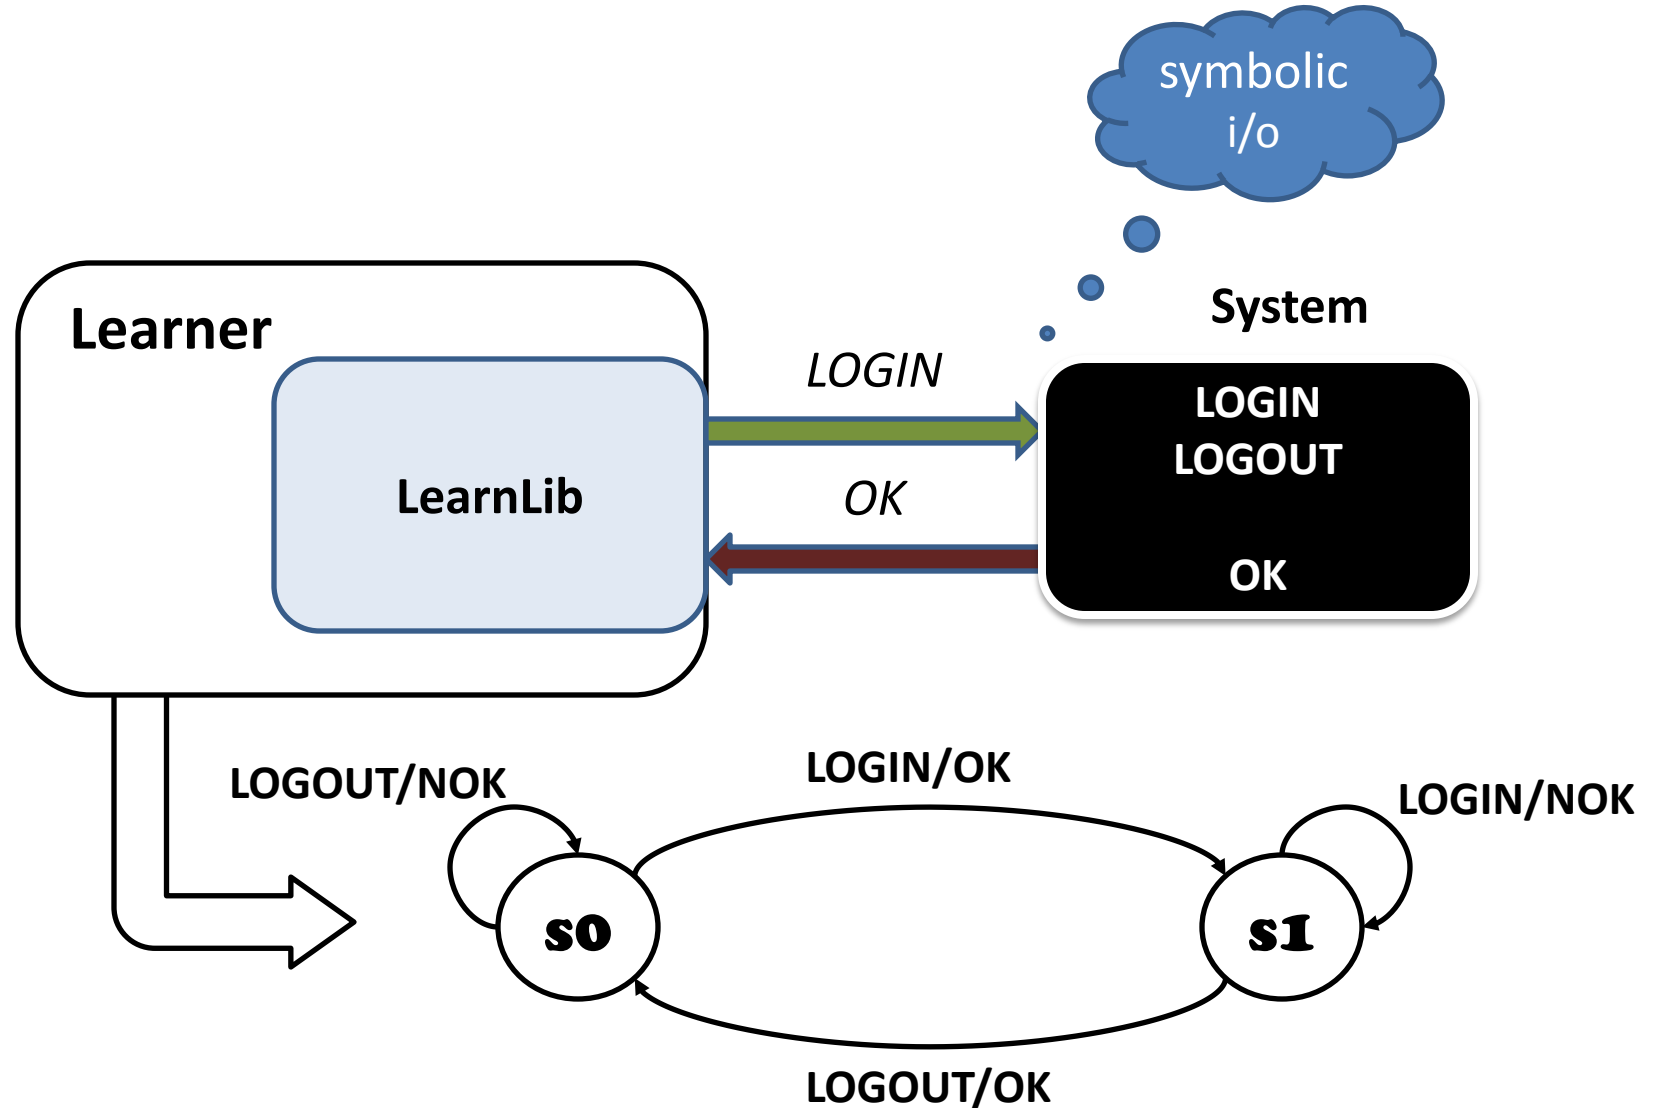
- automatic specification (systems underspecified)
- basis for model based testing + model checking



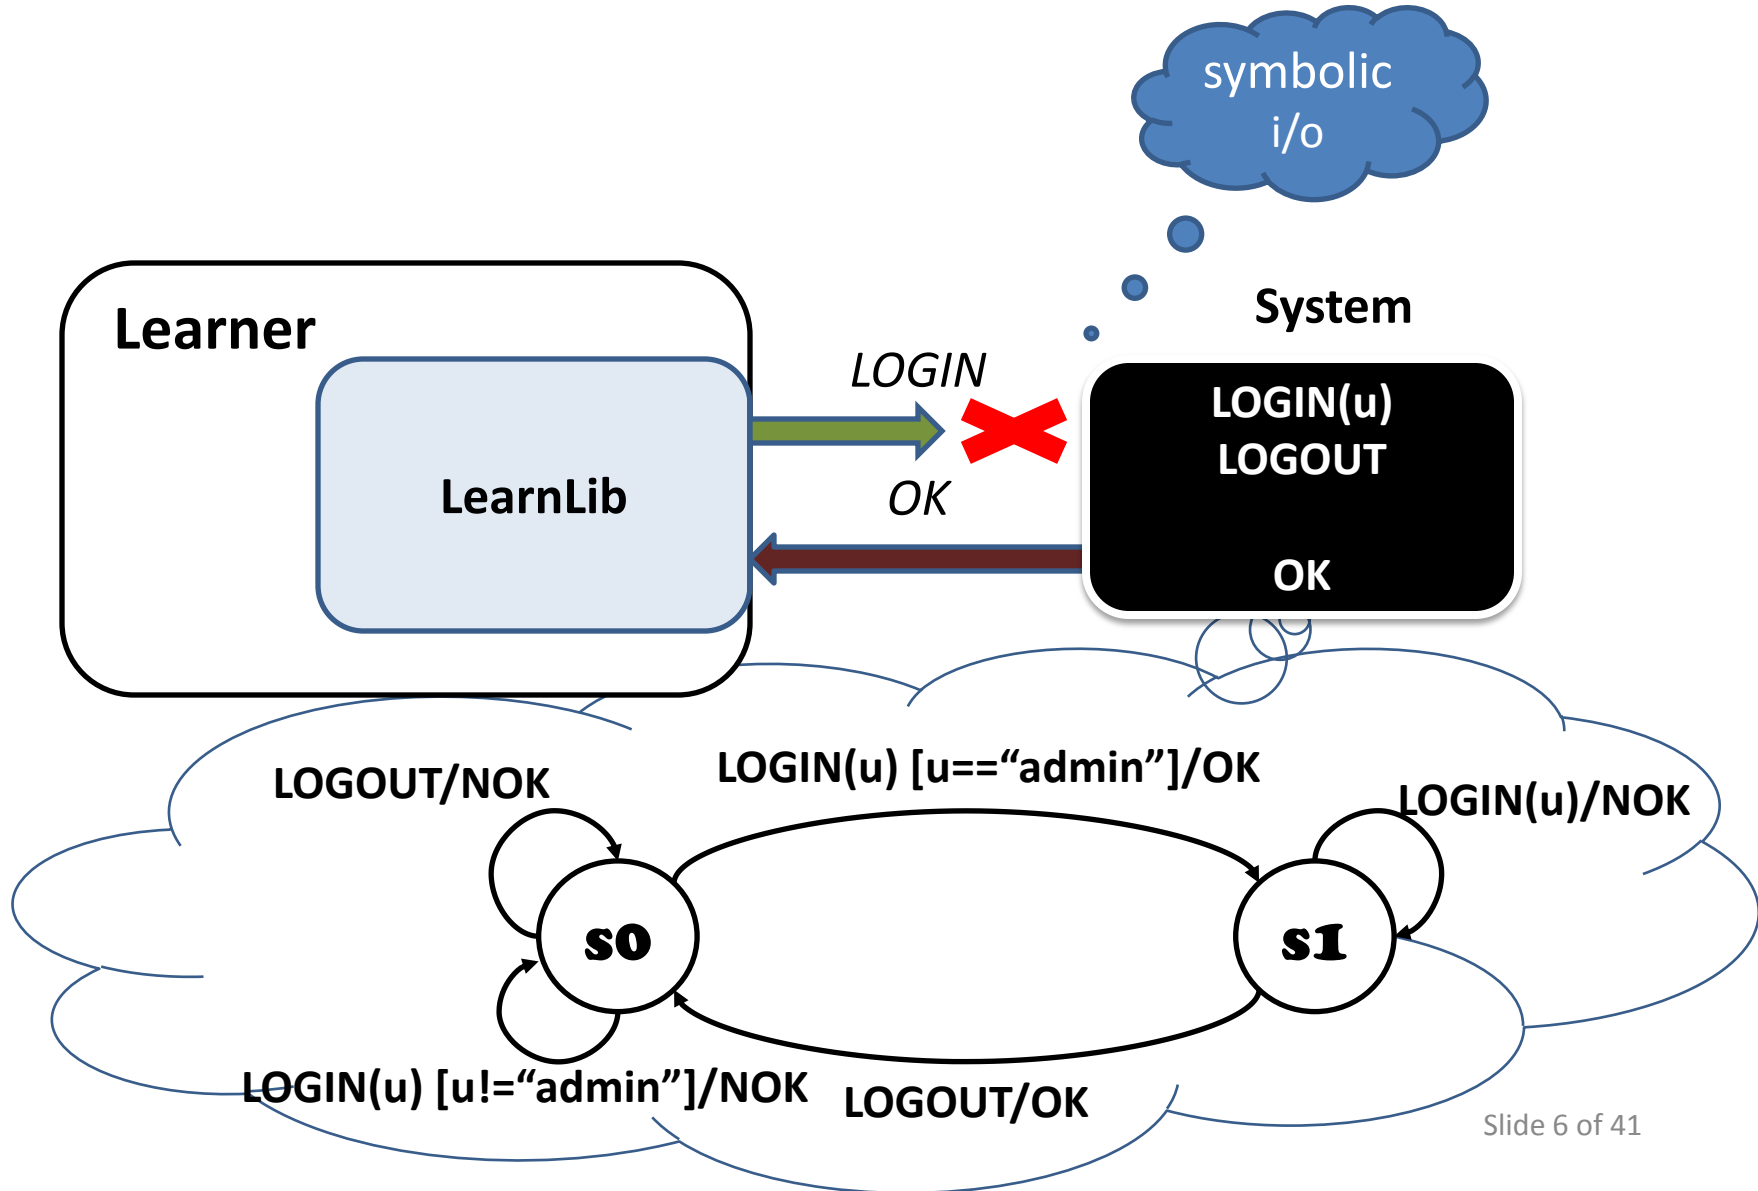
Goals

- apply learning techniques to a realistic system
 - TCP protocol implementations
- show why learning techniques are useful
 - analyze the models obtained against the specification
- establish what we would need to automate
 - milestones for mechanizing the learning process

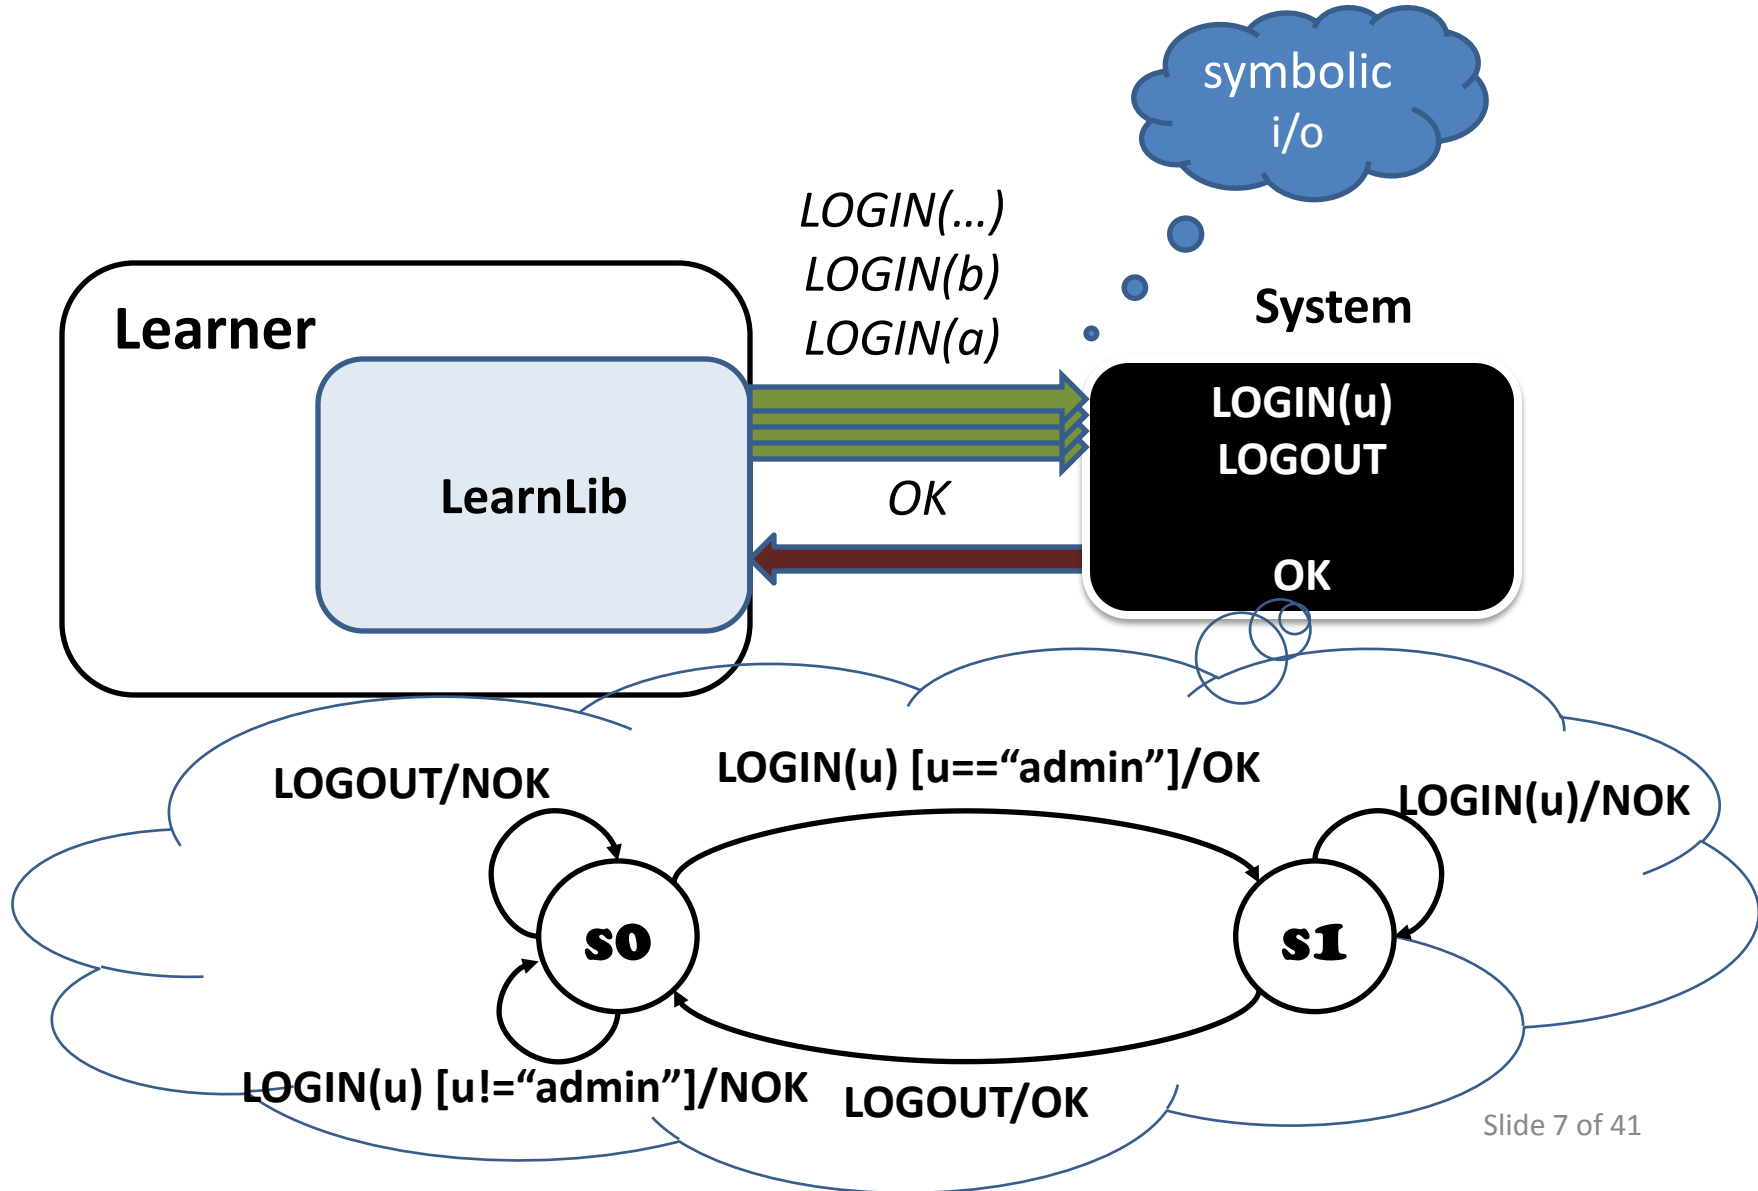
Classical learning



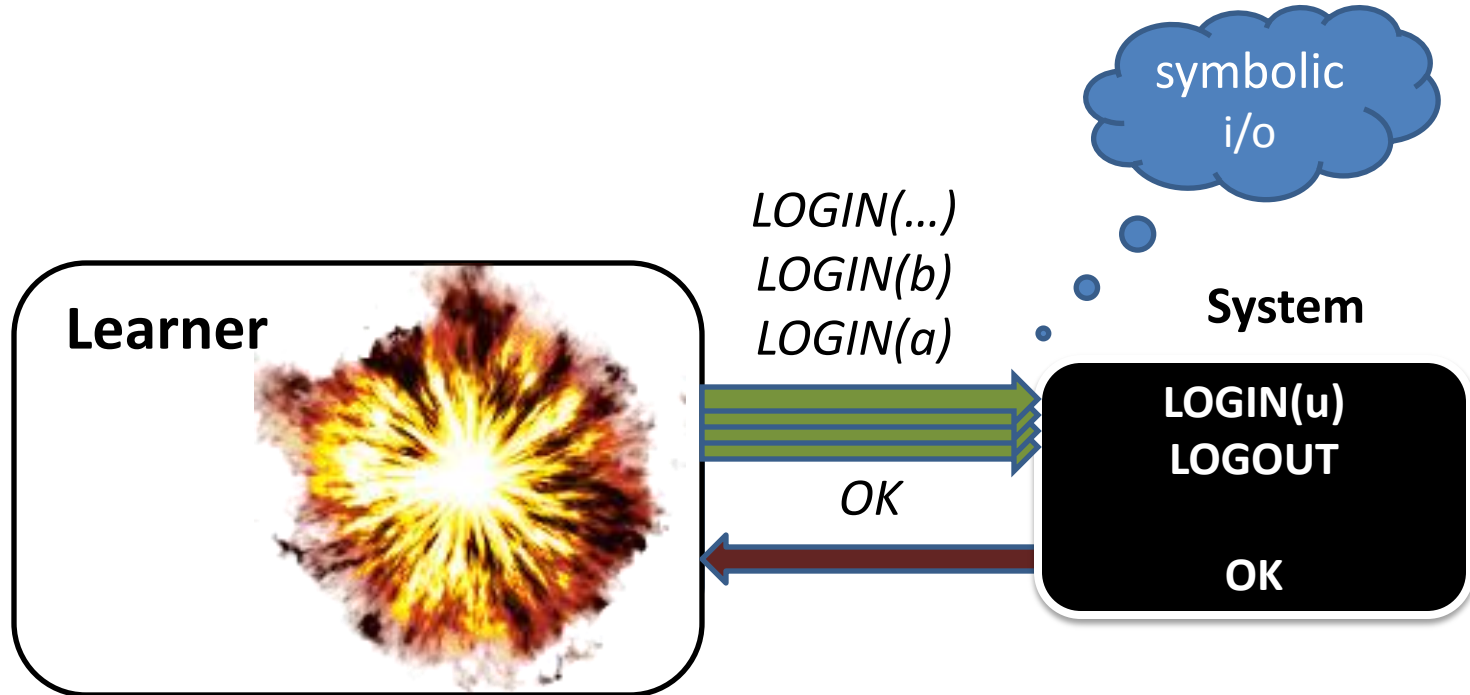
Classical learning



Classical learning

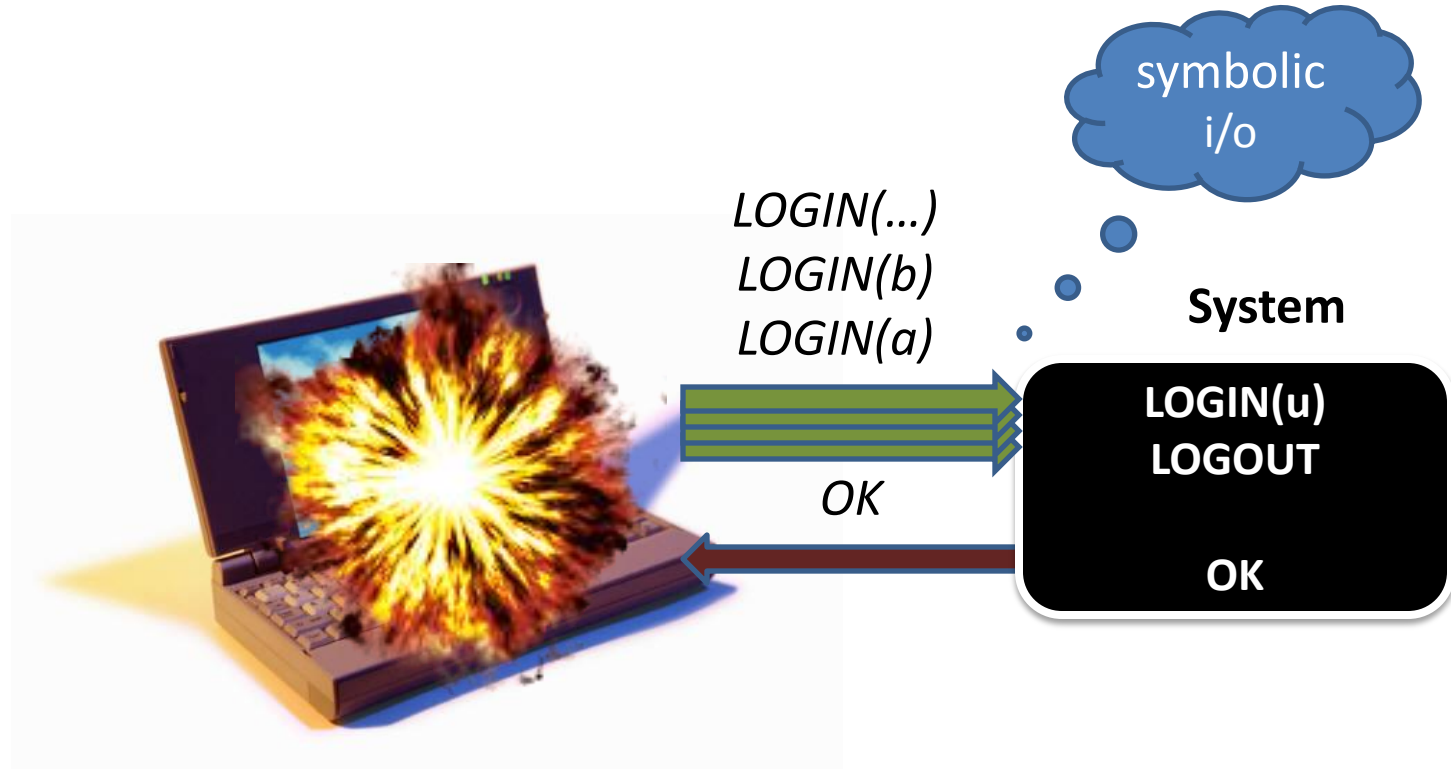


Classical learning



- input symbols for each possible parameter value
- too many for the Learner to handle
 - causes the Learner to explode

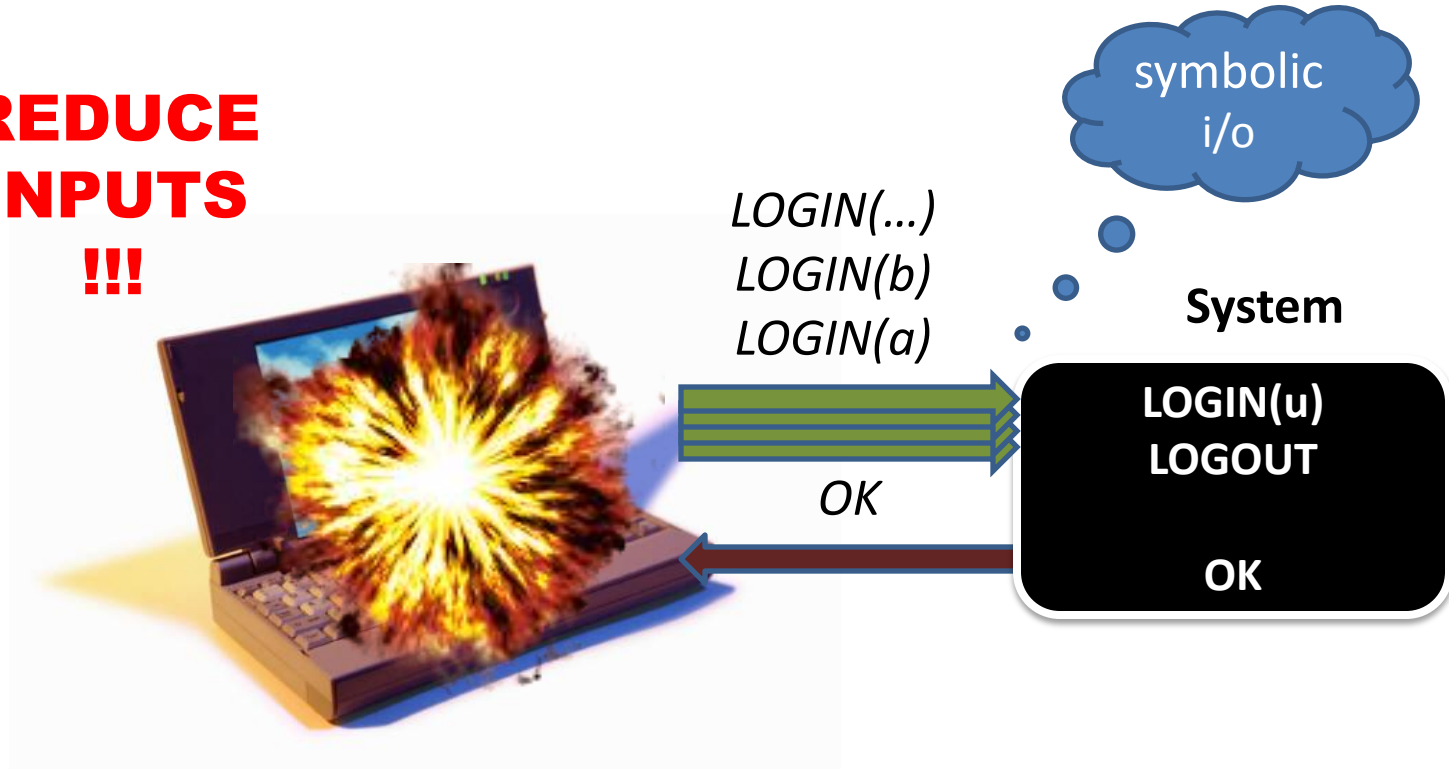
Classical learning



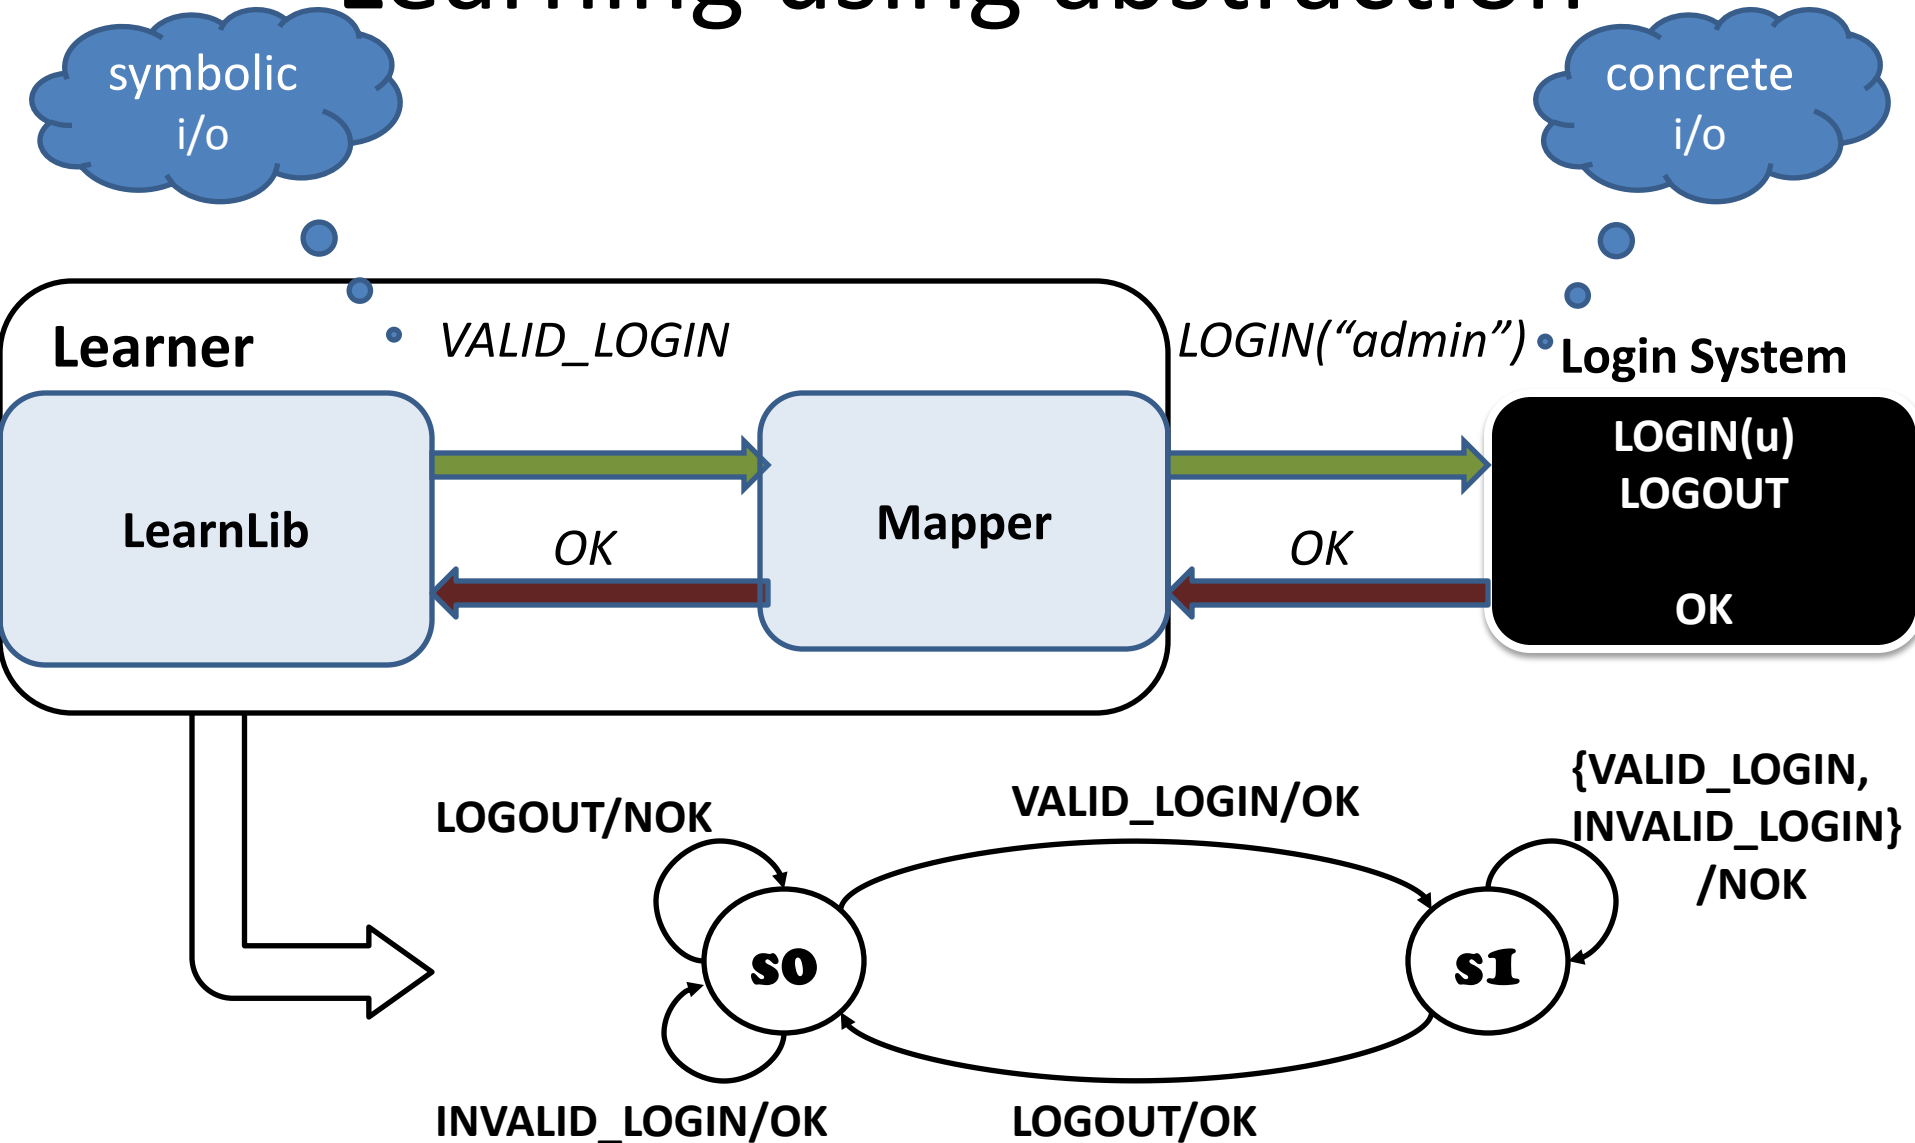
- input symbols for each possible parameter value
- too many for the Learner to handle
 - causes the Learner to explode
 - and the computer

Classical learning

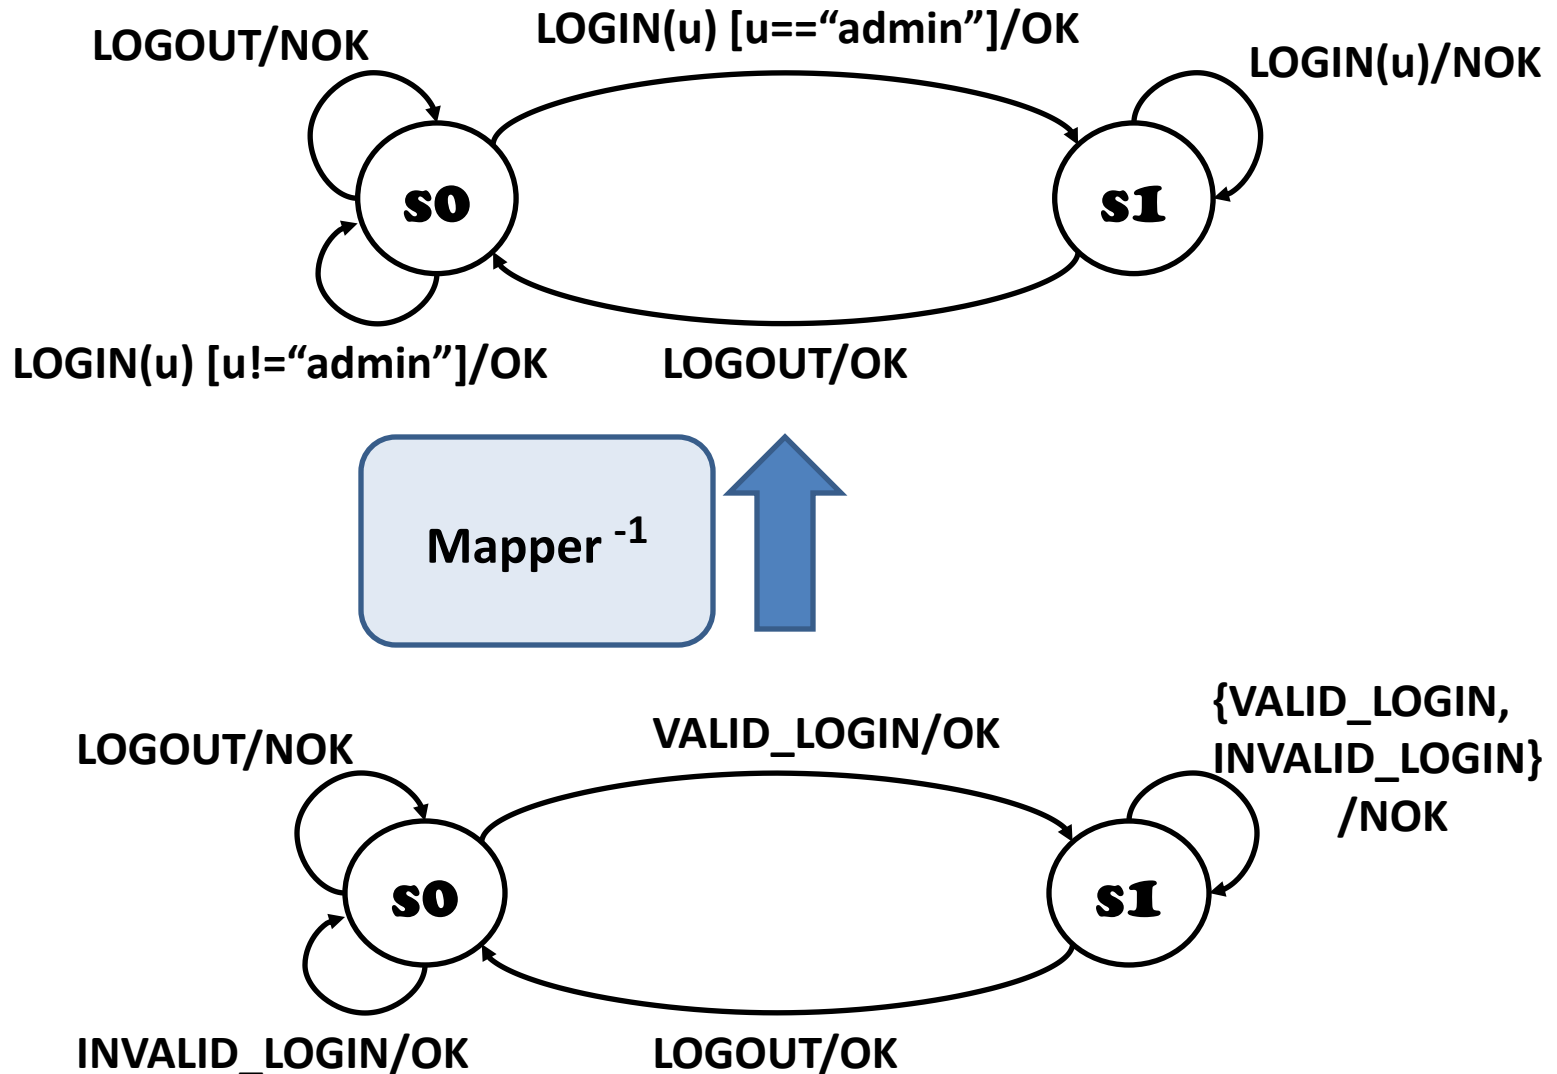
**REDUCE
INPUTS
!!!**



Learning using abstraction

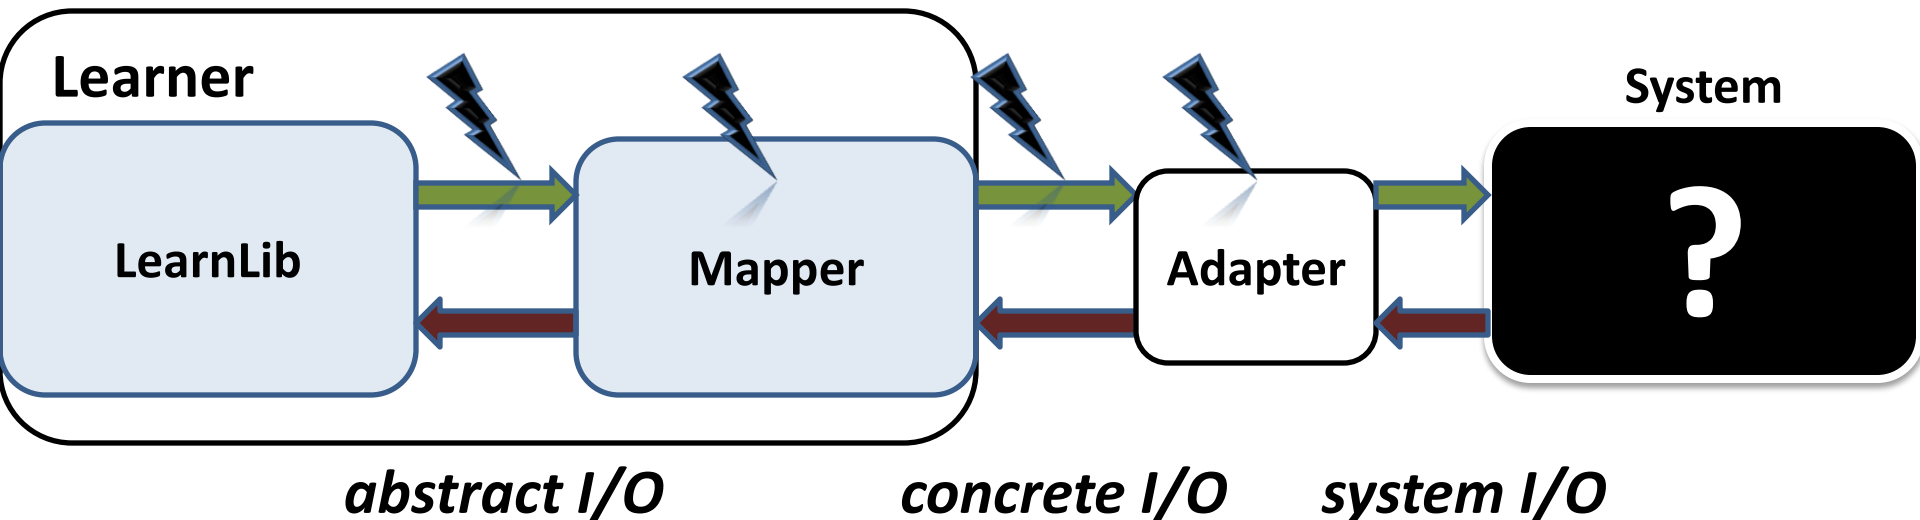


Learning using abstraction



Learning using abstraction

- build **abstract** and **concrete alphabets** (abstract/concrete io)
{VALID_LOGIN, INVALID_LOGIN...} and *{LOGIN(string), LOGOUT...}*
- build **Mapper** component
if input = VALID_LOGIN then return LOGIN("admin")
- build **Adapter** between Learner and System

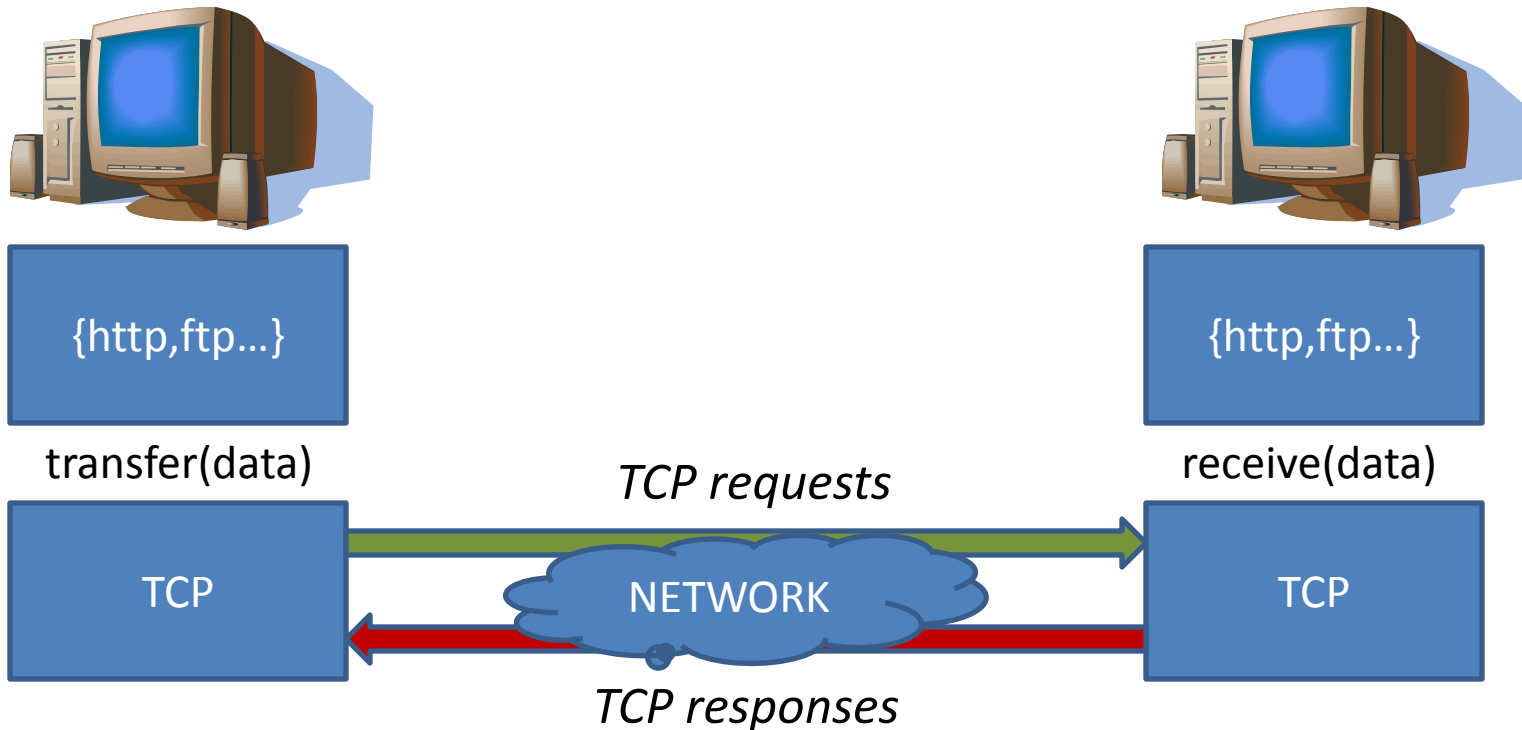


The TCP Protocol

- connection based protocol for reliable data transfer
- connection = (clientIP, clientPort, serverIP, serverPort)

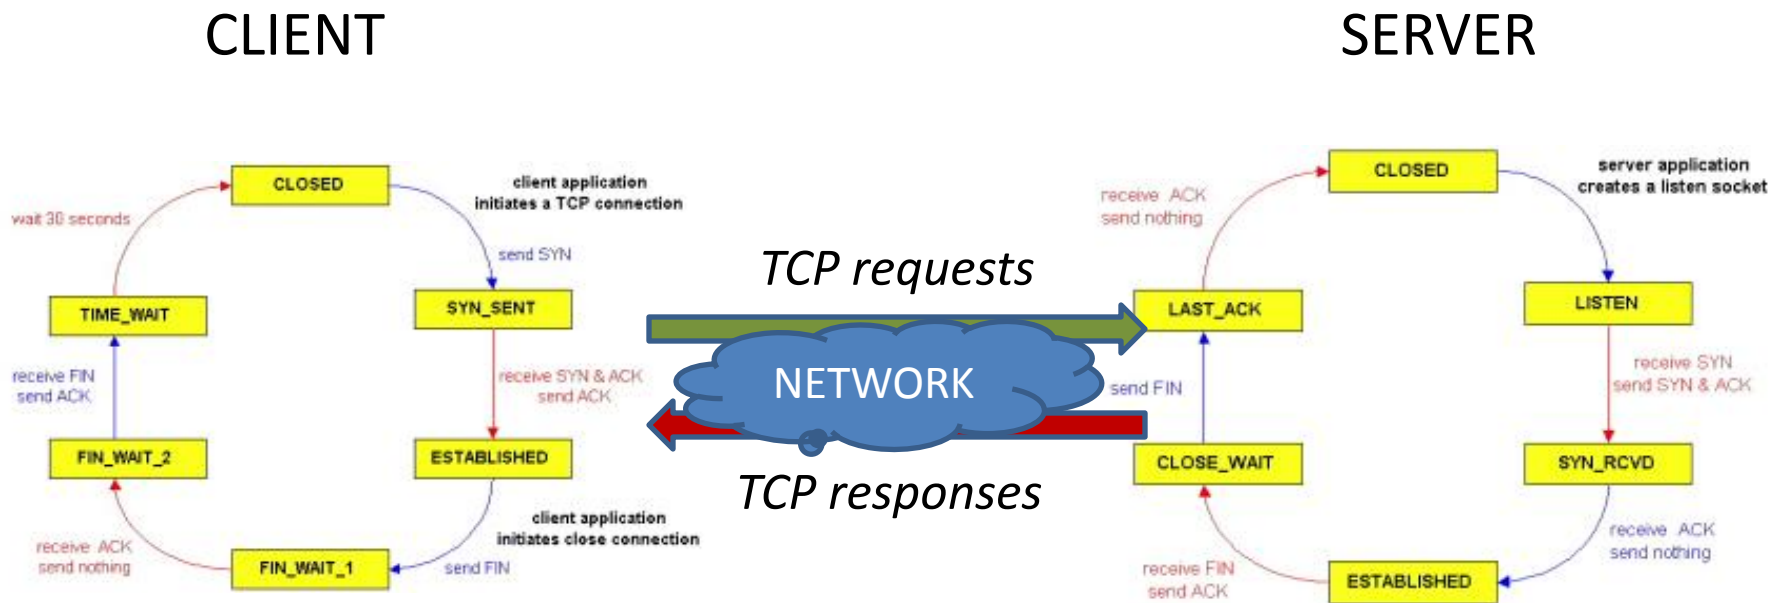
CLIENT APPLICATION

SERVER APPLICATION



Learning the TCP Protocol

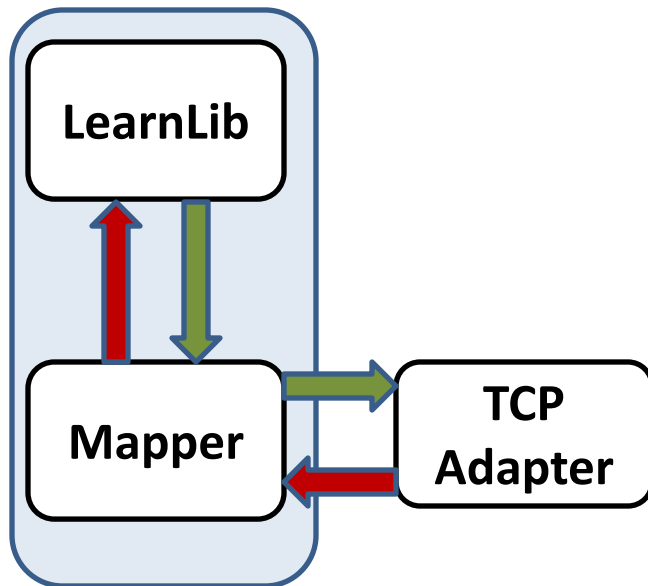
- both the client and the server are state machines



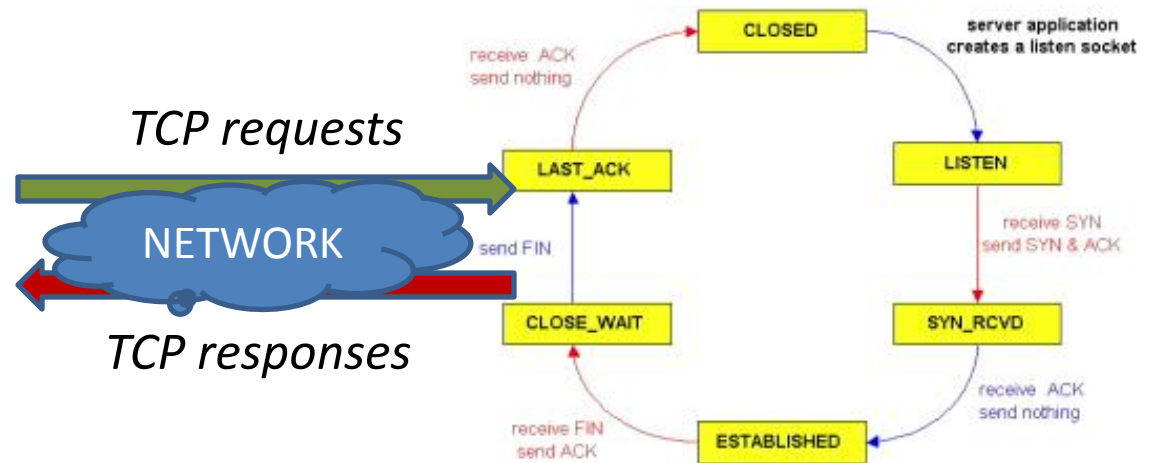
The learning setup

FAKE CLIENT

Learner

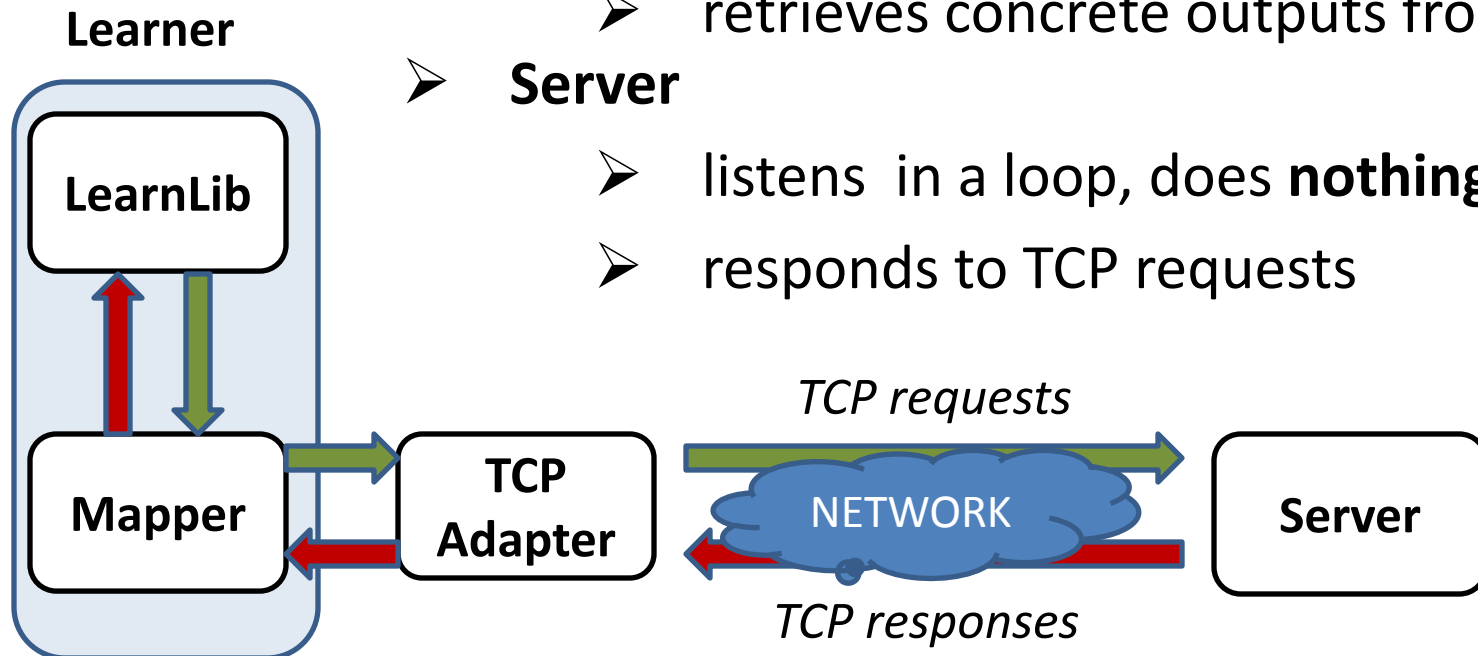


SERVER



The learning setup

- **Learner = LearnLib + Java Mapper**
- **TCP Adapter = Python tool built on Scapy, Pcap, Impacket**
- **Learner**
 - sends concrete inputs to **Adapter**
 - retrieves concrete outputs from **Adapter**
- **Server**
 - listens in a loop, does **nothing** else
 - responds to TCP requests



The alphabet



Abstract Parameters	Values
<i>SeqV, AckV</i>	$\{V, INV\}$
<i>SeqA, AckA</i>	$\{CLSN, SVSN, \dots\}$
<i>Flags</i>	$\{SYN, SYN+ACK, ACK, FIN, FIN+ACK, RST, RST+ACK\}$

Concrete Parameters	Values
<i>SeqNr</i>	$[0, 2^{32}-1]$
<i>AckNr</i>	$[0, 2^{32}-1]$

The mapper

Request(Flags, SeqV, AckV):

if SeqV = V:

SeqNr := clientSeqNr

else:

SeqNr := rand other than clientSeqNr

if AckV = V:

AckNr := serverSeqNr + 1

else:

AckNr := rand other than serverSeqNr

update clientSeqNr

return Request(Flags, SeqNr, AckNr)

abstract to
concrete

update

Request(Flags, SeqV, AckV)



The mapper

Response(**Flags**, SeqNr, AckNr):

SeqA := getAbstractValue(SeqNr)

AckA := getAbstractValue(AckNr)

update clientSeqNr, serverSeqNr

return Response(Flags, SeqA, AckA)

} concrete to
abstract
}
update

getAbstractValue(RegNr):

matches RegNr to clientSeqNr, serverSeqNr,

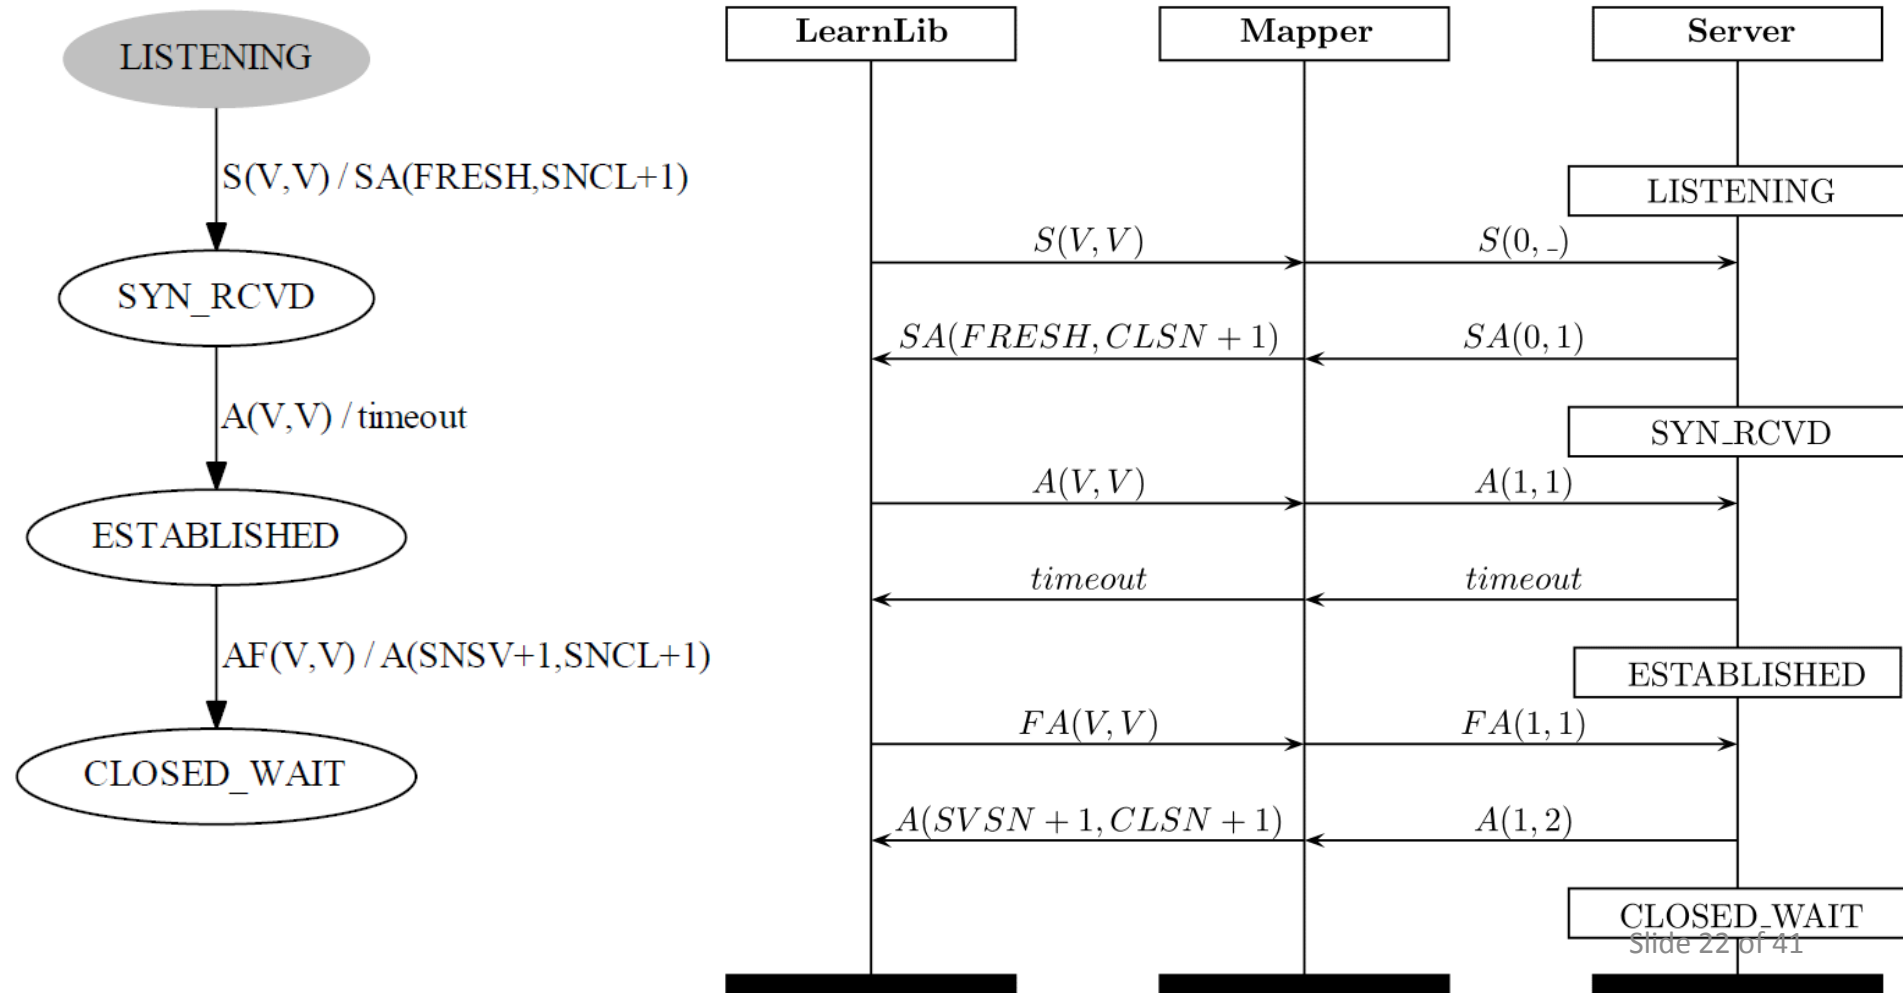
last AckNr./SeqNr. sent/received

results in: CLSN, SVSN, LSS, LAS, ZERO...



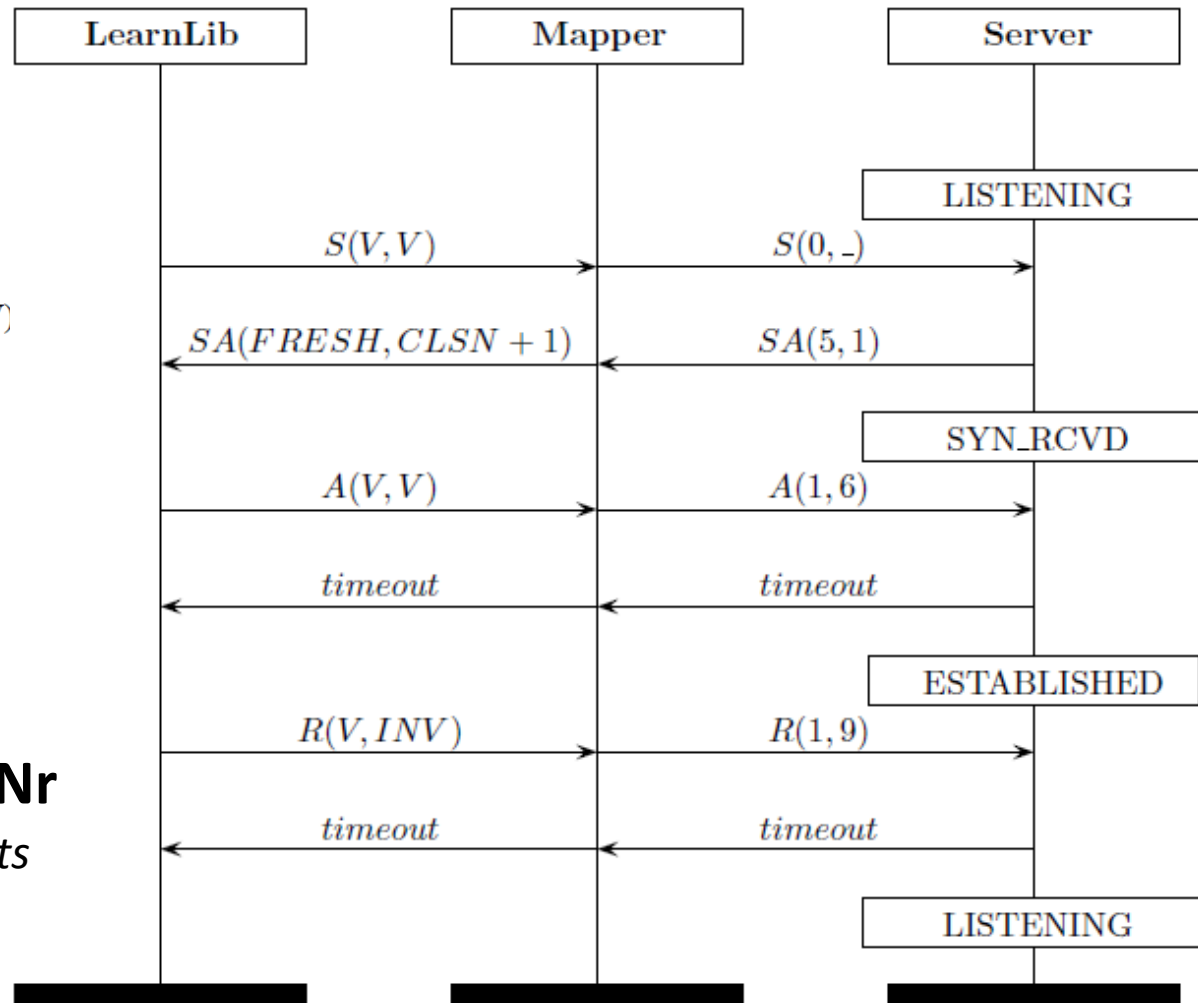
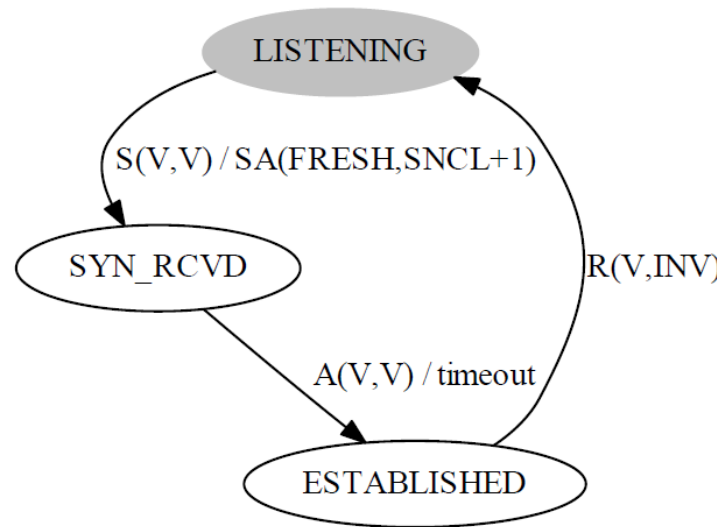
Example trace

TCP Handshake plus closing:



Example trace(2)

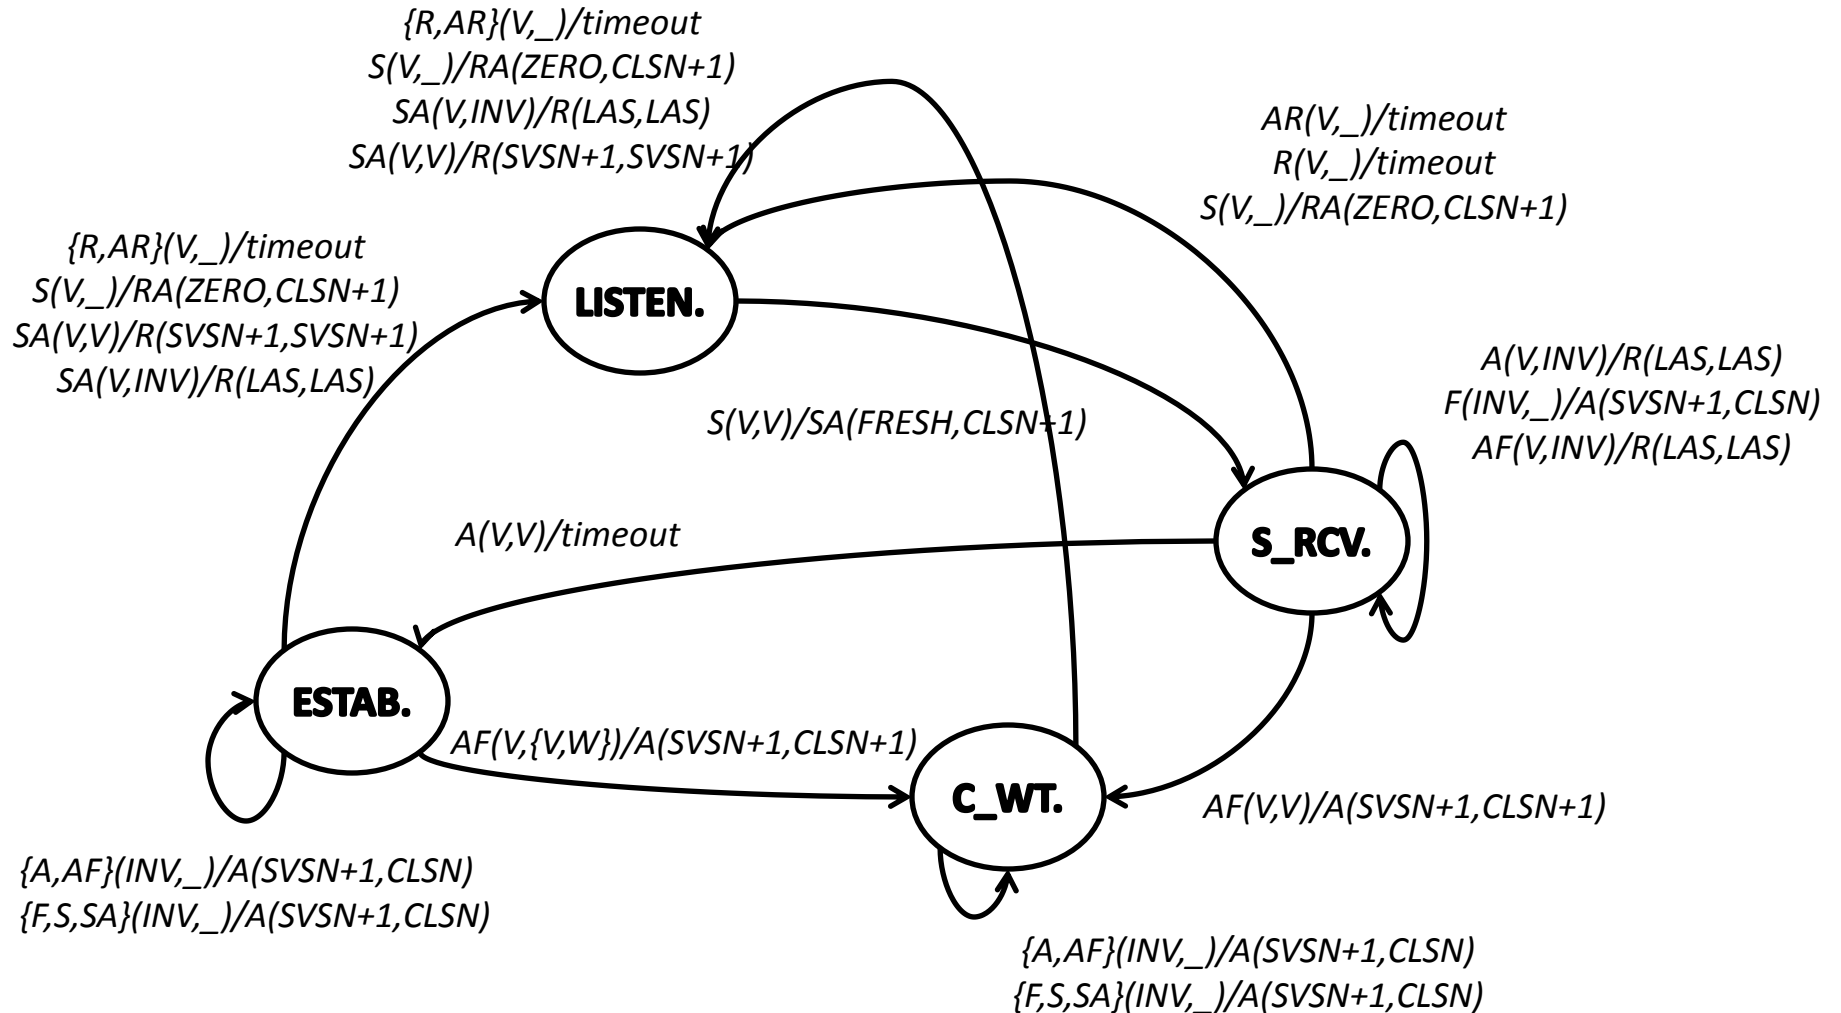
Resetting a connection:



➤ send *RST* with **valid SeqNr**

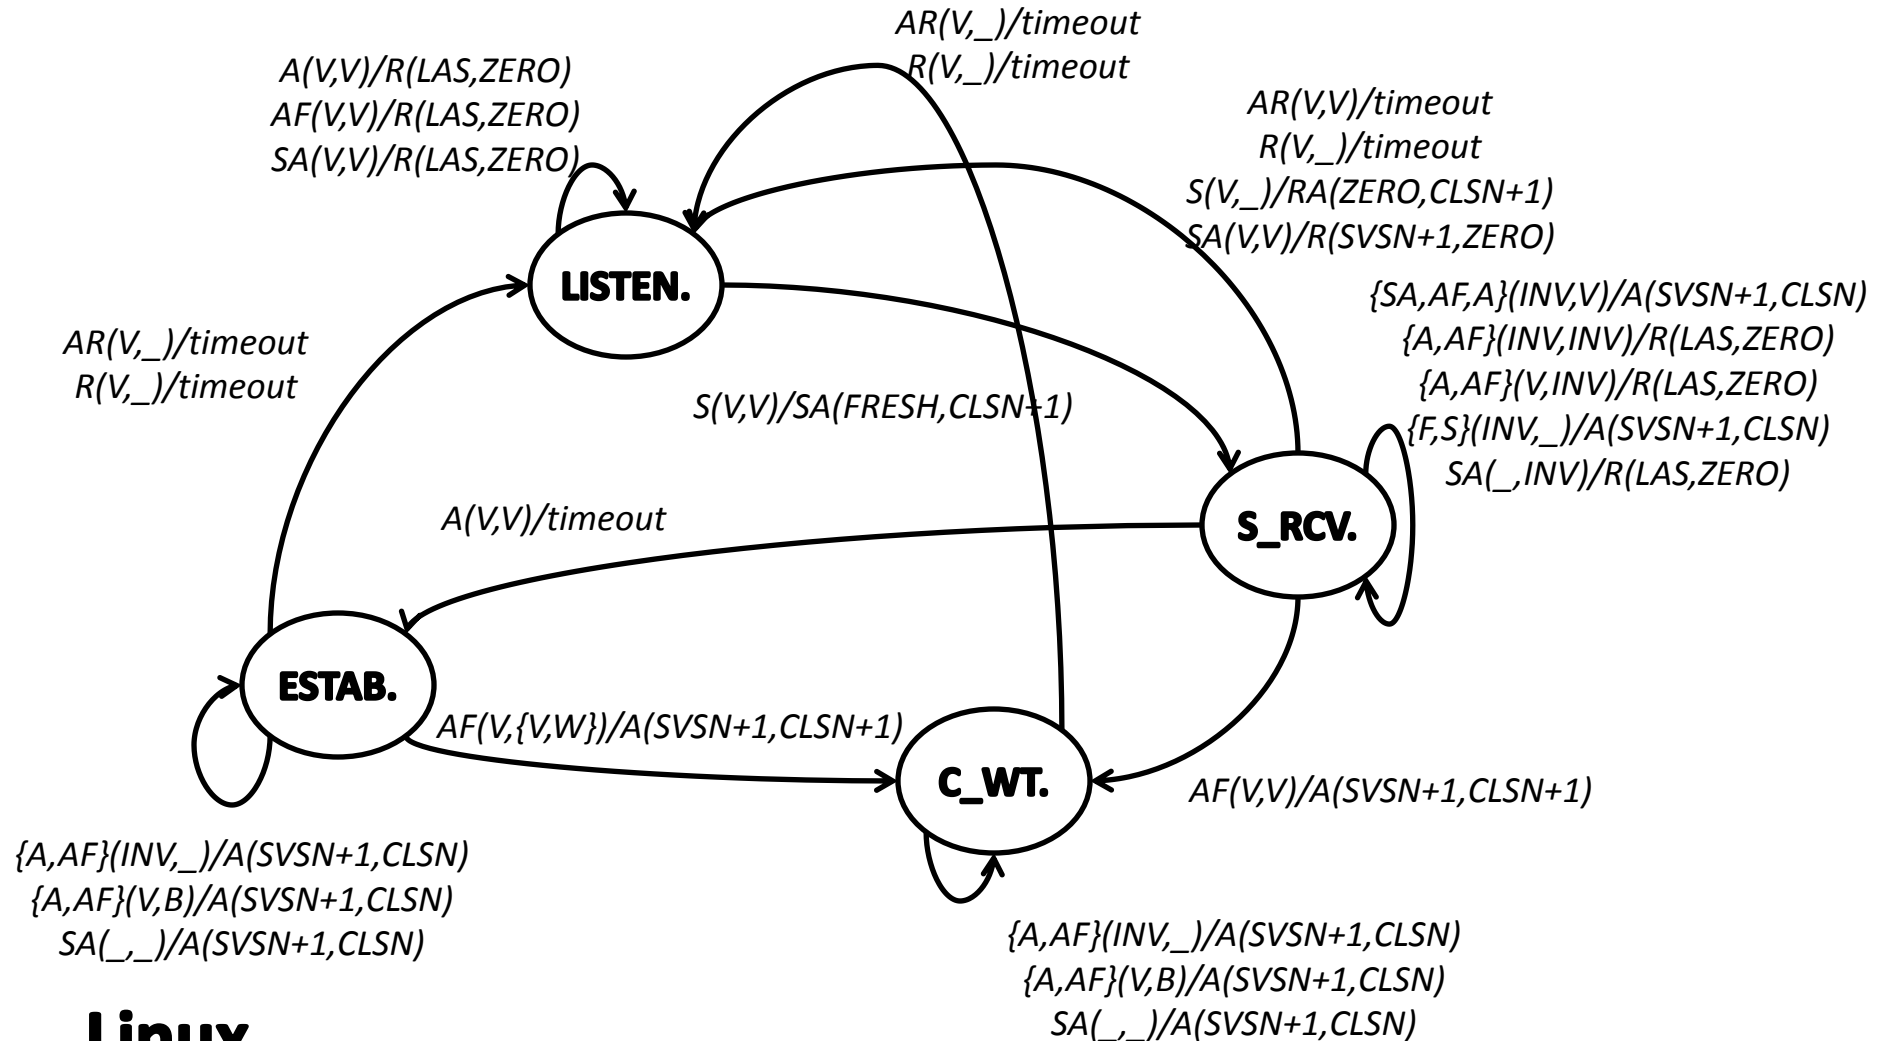
➤ alternative: change ports

Learning results



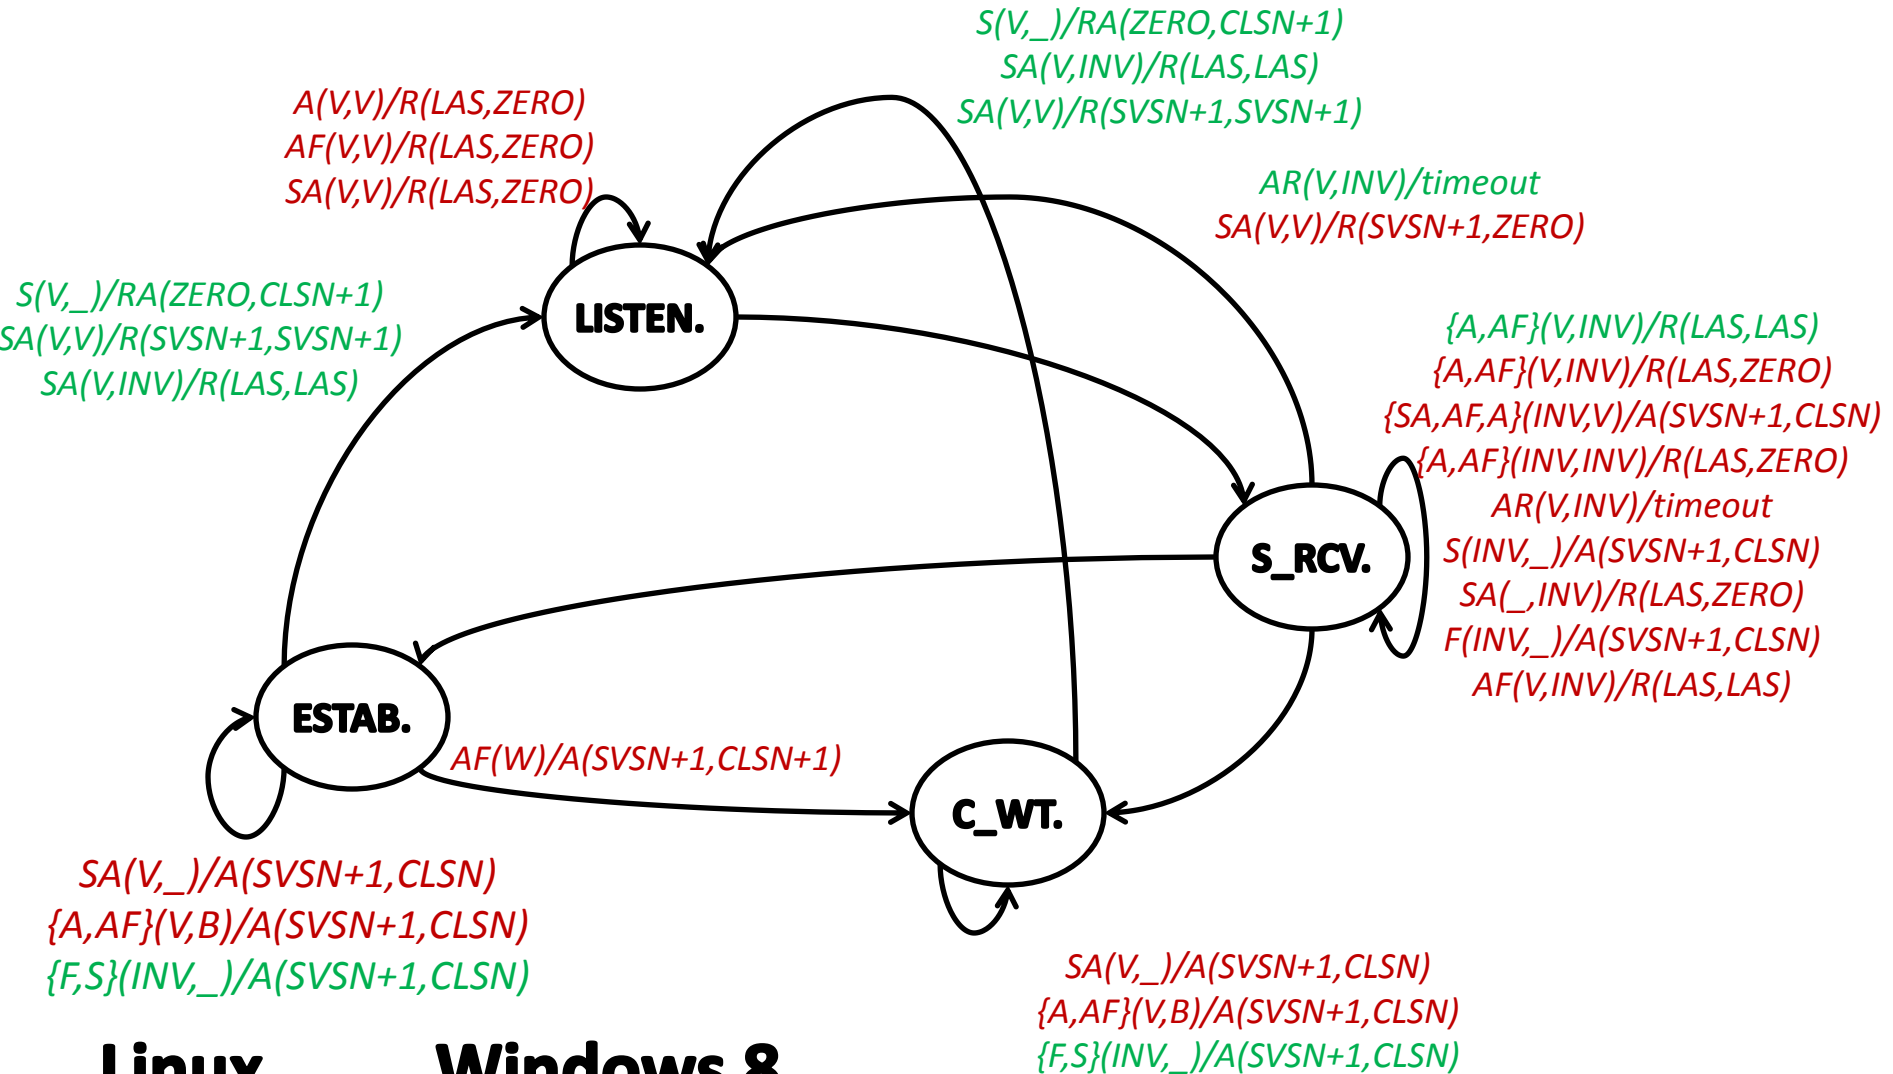
Learning results

$[INV]=[B:W:A]$



Results comparison

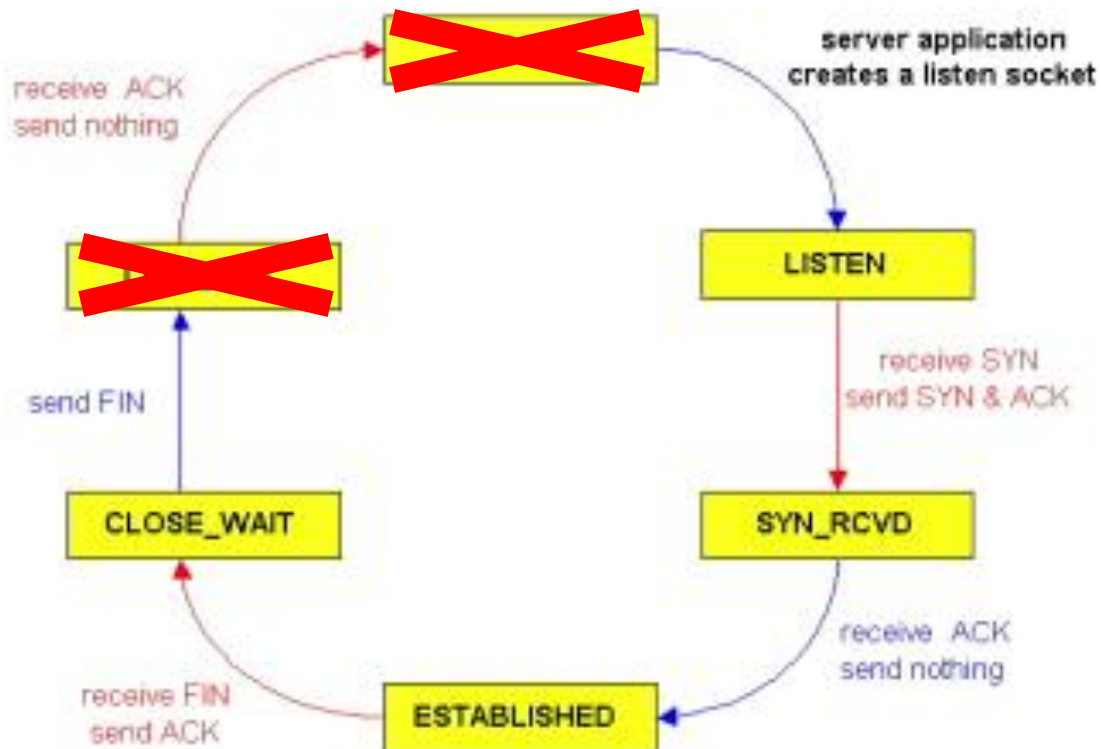
$[INV]=[B:W:A]$



Linux

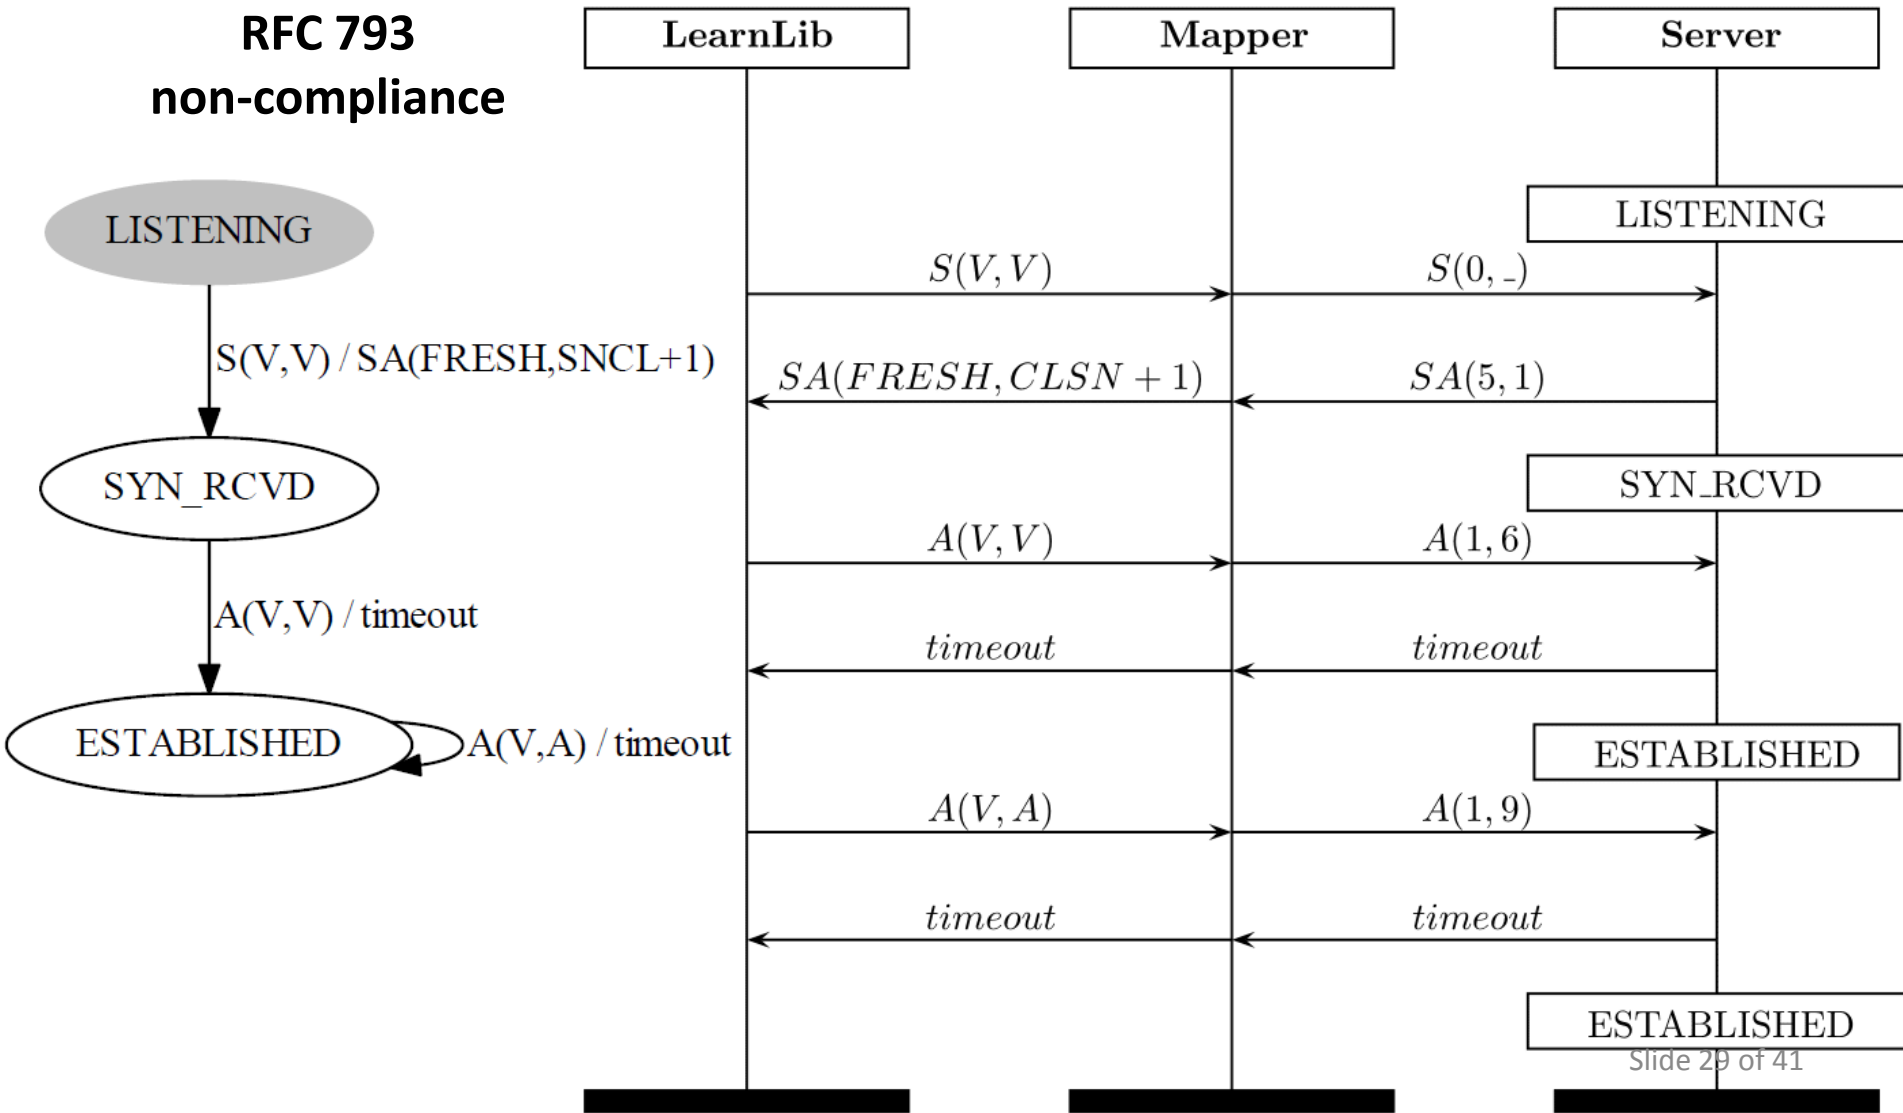
Windows 8

Learning results



Non-compliance

RFC 793 non-compliance



Linux is not compliant

```
/* If the ack includes data we haven't sent yet, discard
 * this segment (RFC793 Section 3.9).
 */
if (after(ack, tp->snd_nxt))
    goto invalid_ack;
...
invalid_ack:
    SOCK_DEBUG (...);
    return -1;
```

Linux is not compliant



If the ACK acks something not yet sent
(...) then **send an ACK**, drop the
segment, and return.

(RFC 793, p. 72)

```
/* If the ack includes data we haven't sent yet, discard  
 * this segment (RFC793 Section 3.9) .  
 */
```

...

Implications

- non-compliance detected
 - specification **not met**, suggested change
- easy comparison between implementations
 - easy OS-fingerprinting
- same approach, same framework to learn other protocols
 - FTP, SCTP...

Future work

- **learn** more TCP

- add more parameters (data, length, window size)
- complicate server (echo server)
- test other implementations

- **learn** higher level protocols

Future work

- **learn** more TCP
 - add more parameters (data, length, window size)
 - complicate server (echo server)
 - test other implementations
- **learn** higher level protocols
 - then the world



Unfortunately...

- writing mappers manually **IS** difficult
 - implies forming abstractions
 - which necessitates reading the specification
 - large number of experiments
 - lengthy (sometimes endless) experiments
 - significant manual labor
 - only rewarding once you learn something



Solution

- Tomte – learning tool
 - can learn a subset of EFSMs **automatically**
 - no mapper required
 - systems learned:
 - login system, SIP protocol, biometric passport
 - more recently: stack and queue systems



Future work



- Tomte – learning tool
 - can learn a **subset** of EFSMs

Tomte	TCP requirements
outputs derived from previous inputs	outputs can be fresh values



Future work

- Tomte – learning tool
 - can learn a **subset** of EFSMs

Tomte	TCP requirements	
outputs derived from previous inputs	outputs can be fresh values	
outputs derived by assignment	outputs derived by simple arithmetic (assignment, increment...)	

Future work

- Tomte – learning tool
 - can learn a **subset** of EFSMs

Tomte	TCP requirements
outputs derived from previous inputs	outputs can be fresh values
outputs derived by assignment	outputs derived by simple arithmetic (assignment, increment...)
time agnostic	time matters



Conclusions

- built a setup for learning a fragment of TCP
 - 4 states, flags, seq and ack numbers, no data
- obtained models which showed functionality
 - and discrepancy Windows/Linux/RFC 793
- set new automation goals
 - so we cut out the middle-man